



Does learning new visual features contribute?



Attribution requires mean ≤ -0.001 and improvement in at least 5/6 seeds. Passing attribution cannot rescue a failed primary screen. The fixed-encoder control still trains its heads and recurrent context/readout.