







Adam | Learning with two memory timescales

Policies run from initialization; no change-triggered reset. Same 16-packet budget, data, and update count.

Uniform history (16) Recent only (16) Recent 8 + historical 8



Lower error and smaller noise penalties are better. Dots show all 6 seeds; later policies have different learned weights.
The fixed split must satisfy acquisition, retention, return, revision, noise, and survival requirements separately for each architecture.